

Numerical Solution of PDEs using the Finite Element Method (www.dealii.org)

Luca Heltai <luca.heltai@sissa.it>

International School for Advanced Studies (www.sissa.it)

Mathematical Analysis, Modeling, and Applications (math.sissa.it)

Master in High Performance Computing (www.mhpc.it)

SISSA mathLab (mathlab.sissa.it)





Main arguments

- Serial scalar poisson solver, in various flavours
- Convergence tests
- Local adaptivity
- Parallelisation techniques in FEM
- If time permits: Vector and Mixed problems, block preconditioning



Tools, Techniques, Best Practices

- What you will learn:
 - Advanced Finite Element theory
 - How to use a modern C++ IDE, to build and debug your codes
 - How to use a large FEM library to solve complex PDE problems
 - How to properly document your code using Doxygen
 - How to use a proper Git workflow to develop your applications
 - How to leverage GitHub actions, google tests, and docker images to test and deploy your application
 - How MPI parallelisation works in real life FEM applications



Outcome of the course

- You will produce your own FEM application based on deal.II which:
 - Solves a PDE of interest to you, on adaptively refined grids, in parallel
 - Uses modern version control tools (on GitHub)
 - Is tested automatically (through GitHub actions) every time you push a commit, or open a pull request
 - Is documented using Doxygen, and its web page is updated and deployed automatically every time you merge to master a new branch



Prerequisites

- Theory:
 - Some knowledge of Sobolev Spaces
 - Linear operators, Banach and Hilbert spaces, duality, etc.
 - One elementary course on Numerical Analysis
 - Quadrature, interpolation, Taylor expansions, etc.
- Practice (for the first few lectures):
 - a machine with Visual Studio Code installed (c++-11 is required)
 - Docker
 - A GitHub account



More Info

- Course pages:

- Course main page, with schedule and up to date information

<https://www.math.sissa.it/course/phd-course/fem-with-dealii-2022>

- Course slides, notes, materials, and codes:

<https://github.com/dealii-courses/fem-with-dealii-2022>

- Email:

prof. Luca Heltai <luca.heltai@sissa.it>

Numerical Solution of PDEs using the Finite Element Method

Setting up a “best practice” FEM development environment
(some slides taken from prof. Wolfgang Bangerth)

Luca Heltai <luca.heltai@sissa.it>

International School for Advanced Studies (www.sissa.it)

Mathematical Analysis, Modeling, and Applications (math.sissa.it)

Master in High Performance Computing (www.mhpc.it)

SISSA mathLab (mathlab.sissa.it)



Why deal.II (or any other Finite Element library)

- The numerical solution of partial differential equations is an immensely vast field!
- It requires us to know about:
 - Partial differential equations
 - Methods for discretizations, solvers, preconditioners
 - Programming
 - Adequate tools
- **This course will cover all of this to some degree!**



Numerics of PDEs

There are 3 standard tools for the numerical solution of PDEs:

Finite element method (FEM)

Finite volume method (FVM)

Finite difference method (FDM)

Common features:

Split the domain into small volumes (cells)

Define balance relations on each cell

Obtain and solve very large (non-)linear systems

Problems:

Every code has to implement these steps

There is only so much time in a day

There is only so much expertise anyone can have



Numerics of PDEs

There are 3 standard tools for the numerical solution of PDEs:

- Finite element method (FEM)
- Finite volume method (FVM)
- Finite difference method (FDM)

Common features:

- Split the domain into small volumes (cells)
- Define balance relations on each cell
- Obtain and solve very large (non-)linear systems

In addition:

- We don't just want a simple algorithm
- We want state-of-the-art methods for everything



Numerics of PDEs

Examples of what we would like to have:

Adaptive meshes

Realistic, complex geometries

Quadratic or even higher order elements

Multigrid solvers

Scalability to 1000s of processors

Efficient use of current hardware

Graphical output suitable for high quality rendering

Q: How can we make all of this happen in a single code?



The hard reality

- Most research software today:
 - **Written by graduate students**
 - without a good overview of existing software
 - with little software experience
 - with little incentive to write high quality code
 - **Maintained by postdocs**
 - with little time
 - who need to consider the software primarily as a tool to publish papers
 - **Advised by faculty**
 - with no time
 - oftentimes also with little software experience



How we develop Software

Q: How can we make all of this happen in a single code?

Not a question of feasibility but of how we develop software:

Is every student developing their own software?

Or are we re-using what others have done?

Do we insist on implementing everything from scratch?

Or do we build our software on existing libraries?



How we develop Software

Q: How can we make all of this happen in a single code?

Not a question of feasibility but of how we develop software:

Is every student developing their own software?

Or are we re-using what others have done?

Do we insist on implementing everything from scratch?

Or do we build our software on existing libraries?

There has been a major shift on how we approach the second question in scientific computing over the past 10-15 years!



**The secret to good scientific software is
(re)using existing libraries!**



Existing Software

There is excellent software for almost every purpose!

Basic linear algebra (dense vectors, matrices):

BLAS
LAPACK

Parallel linear algebra (vectors, sparse matrices, solvers):

PETSc
Trilinos

Meshes, finite elements, etc:
deal.II – the topic of this class

...

Visualization, dealing with parameter files, ...



Our experience

It is realistic for a student developing numerical methods using external libraries to have a code at the end of a PhD time that:

- Works in 2d and 3d
- On complex geometries
- Uses higher order finite element methods
- Uses multigrid solvers or preconditioners
- Solves a nonlinear, time dependent problem

Doing this from scratch would take 10+ years.



Common arguments...

Arguments against using other people's packages:

I would need to learn a new piece of software, how it works, its conventions. I would have to find my way around its documentation. Etc.
I think I'll be faster writing the code I want myself!



Common arguments...

Arguments against using other people's packages:

I would need to learn a new piece of software, how it works, its conventions. I would have to find my way around its documentation. Etc.
I think I'll be faster writing the code I want myself!

Answers:


The first part is true.

The second is not!

You get to use a lot of functionality you could never in a lifetime implement yourself.
Think of how we use Matlab today!




I'm faster!



Blue: use external libraries
Red: do it yourself



The real picture...



Blue: use external libraries
Red: do it yourself



Common Arguments...

Arguments against using other people's packages:

I want my students to actually understand the methods they are doing. So I let them code things from scratch!



Common Arguments...

Arguments against using other people's packages:

I want my students to actually understand the methods they are doing. So I let them code things from scratch!

Answers:

Yes, there is value to that.

But: if you know quadrature in 2d, why implement it again in 3d?

So let them write a toy code and throw it away after 3 months and do it right based on existing software.



Common Arguments...

Arguments against using other people's packages:

How do I know that that software I'm supposed to use doesn't have bugs? How can I *trust* other people's software?

With my own software, at least I know that I don't have bugs!

Answer 1:

You can't be serious to think that your own software has no bugs!



Common Arguments...

Arguments against using other people's packages:

How do I know that that software I'm supposed to use doesn't have bugs? How can I *trust* other people's software?

With my own software, at least I know that I don't have bugs!

Answer 2:

The packages I will talk about are developed by professionals with a lot of experience

They have extensive testsuites

For example, deal.II runs 3,000+ tests **after every single change**



Bottomline:

When having to implement software for a particular problem, re-use what others have done already

There are many high-quality, open source software libraries for every purpose in scientific computing

Use them:

- You will be far more **productive**
- You will be able to use **state-of-the-art** methods
- You will have far **fewer bugs** in your code

If you are a graduate student:

Use them because you will be able to impress your advisor with quick results!



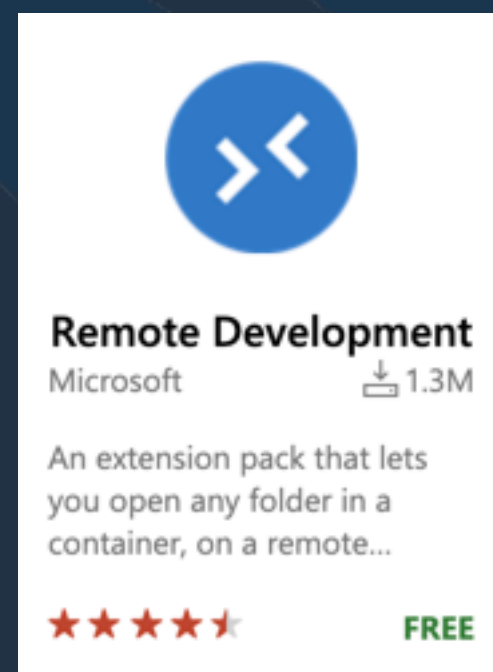
Roadmap for next lectures:

- Version control system (git)
- Modern IDE (VSCode)
- Cross platform build systems (cmake)
- Automatic formatting (clang-format)
- Test driven development (google test, deal.II testing framework)
- Inline documentation (doxygen)
- External visualisation tools (paraview)



Setting up VSCode

- Download and install **Docker**: <https://www.docker.com/products/docker-desktop>
 - Read some doc: <https://www.docker.com/get-started>
- Download and install: <https://code.visualstudio.com/download>
 - Read some doc: <https://code.visualstudio.com/docs>
 - Install the following extension:



Open the directory of the repository

- Open the directory containing the repository of the course. The directory contains a hidden folder, called “**.devcontainer**”
- VSCode should ask you if you want to run the folder in the container. Say yes.
- VSCode will now download a docker image. The first time around, this will take some time.